

AI Hosting & Cost Report

KEYSS Timesheet Assistant | Model: gwen3-30b | Prepared for management review

Prices approximate, INR at Rs 85 = \$1. API rates web-verified Jun 2026; GPU rates from providers' live pages.

RECOMMENDED — Cloudflare Workers AI is the BEST option for us

Cheapest (~Rs 500/month for ~20 users), keeps employee data private (no training), zero setup. Slightly slower (3-4s) but reliable.

1. How the cost is calculated (step by step)

Every option's cost depends on how much the AI is used. We build it up from one request, to one user, to the whole team.

Step	Value
Input tokens per request	5,000
Output tokens per request	1,000
Total tokens per request	6,000
Requests per user per day	20
Tokens per user per day	120,000
Tokens per user per month (x30)	3,600,000 (3.6M)
Cost per user per month (Cloudflare: $3.0M \times \$0.051 + 0.6M \times \0.34)	~\$0.36 = Rs 30

So on Cloudflare each user costs about **Rs 30/month**. Total = Rs 30 x number of users (minus the free tier).

Total cost on Cloudflare by team size (after the free 10,000 neurons/day):

Users	Requests/day	Tokens/month	Cost/month
10	200	36M	~Rs 0-150 (mostly free)
20	400	72M	~Rs 400-500
50	1,000	180M	~Rs 1,200
100	2,000	360M	~Rs 2,700

Cloudflare gives the first ~185-270 requests/day FREE, which is why a small team (10 users) pays almost nothing. Cost grows in a straight line with usage. The same method applies to other providers using their own rates (Section 5).

2. Two ways to run AI

(a) Ready API - a company hosts the model; we pay per use, no server to manage (Cloudflare, DeepInfra, OpenAI...).

(b) Rent a GPU + run our own - rent a graphics server (RunPod, Vast.ai) and install free software (**Ollama** / **vLLM**) on it. Note: Ollama/vLLM are FREE software, not a service - their cost is the machine they run on, and they need a GPU (a CPU server/laptop times out).

3. Full comparison

Service	Type	Speed	Cost @ 20 users	Private?	Verdict
Cloudflare Workers AI	Ready API	Slow (3-4s)	~Rs 400-600/mo	Yes	BEST - cheapest, no setup
DeepInfra	Ready API	Fast	~Rs 1,800/mo	Yes	Cheap + fast
Together.ai	Ready API	Fast	~Rs 1,700-5,000/mo	Yes	Good, costlier
Fireworks.ai	Ready API	Fast	~Rs 1,700-5,000/mo	Yes	Similar to Together
OpenAI GPT-4o-mini	Ready API	Fast	~Rs 1,400/mo	Yes	Very reliable
OpenAI GPT-4o (full)	Ready API	Fast	~Rs 23,000/mo	Yes	Top quality, expensive
Anthropic Claude Haiku 4.5	Ready API	Fast	~Rs 6,000-10,000/mo	Yes	Best tool-calling + caching
Google Gemini 2.0 Flash	Ready API	Fast	~Rs 900/mo (paid)	Free=No / Paid=Yes	Free tier TRAINS data

Service	Type	Speed	Cost @ 20 users	Private?	Verdict
OpenRouter Nemotron (free)	Ready API	Medium	Rs 0	No (trains)	AVOID - privacy risk
OpenRouter premium models	Ready API	Fast	~Rs 98,800/mo	Yes	Too expensive at scale
Groq	Ready API	Fastest (1.1s)	tight limits	Yes	Limits choke a team
RunPod Serverless (RTX 4090)	Rent GPU	Fast (warm)	~Rs 1,600-2,000/mo	Yes	Fast + private, pay/use
Vast.ai RTX 4090 (24/7)	Rent GPU	Fast	~Rs 22,000/mo	Yes	Costly if always on
Vast.ai RTX 5090 32GB (24/7)	Rent GPU	Very fast	~Rs 27,000/mo	Yes	Full quality, always on
Ollama / vLLM	Free software	= the GPU	= GPU cost	Yes	Install on rented GPU
DigitalOcean / Hetzner	CPU server	Unusable	Rs 1,300-3,000/mo	-	NO GPU - cannot run AI

4. Privacy - safe for company data?

Safe (data NOT used for training): Cloudflare, DeepInfra, Together, Fireworks, OpenAI (API), Claude (API), Gemini PAID, and any GPU we rent ourselves (RunPod/Vast + Ollama/vLLM).

NOT safe (logs / trains on our data): OpenRouter FREE models (e.g. Nemotron free) - fully free but their policy trains on submitted data; and Google / Gemini FREE tier also trains on data. Our timesheet data has employee info, so these should be used only for testing with dummy data, not real data.

5. Monthly cost at 20 users (~72M tokens/month)

Option	Rate used	Monthly cost
Cloudflare	\$0.051/M in, \$0.34/M out (minus free)	~Rs 400-600
DeepInfra	\$0.20/M in, \$0.80/M out	~Rs 1,800
Gemini 2.0 Flash (paid)	\$0.10/M in, \$0.40/M out	~Rs 900
OpenAI GPT-4o-mini	\$0.15/M in, \$0.60/M out	~Rs 1,400
RunPod serverless (4090)	\$0.00031/sec (400-500 calls/day)	~Rs 1,600-2,000
Claude Haiku 4.5	\$1/M in, \$5/M out (caching helps)	~Rs 6,000-10,000
OpenAI GPT-4o (full)	\$2.50/M in, \$10/M out	~Rs 23,000
OpenRouter premium (per request)	~\$0.0968 per request x 12,000	~Rs 98,800 (very expensive)
Vast.ai 4090 / 5090 (24/7)	\$0.37 / \$0.45 per hour	Rs 22,000 / Rs 27,000
OpenRouter Nemotron (free)	\$0	Rs 0 - but TRAINS data

RunPod is charged by GPU-time, so cost depends on calls/day: ~Rs 1,600 at 400/day, ~Rs 2,000 at 500/day, ~Rs 3,900 (\$45.83) at 1,000/day. The per-request premium models (~\$0.0968/request) reach ~Rs 1 lakh/month at our scale.

6. How many users each option supports

Option	Comfortable users	Notes
Cloudflare (free)	~9-13 users	Free ~185-270 calls/day
Cloudflare (paid)	20-100+	Auto-scales, very cheap per call
DeepInfra / OpenAI / Together	20-500+	Scale freely, pay per use
RunPod serverless (4090)	20-100+	Auto-scales GPU workers (vLLM)
One rented GPU (always-on)	20-80 concurrent	Handles 5-6+ at once, no crash
Groq (free)	1-2 users only	8K tokens/min blocks a team

Will it crash if 5-6 people use it together? No. Ready APIs and GPU servers (with vLLM) handle many simultaneous users by queuing/batching/scaling. Only Groq's per-minute limit is a problem for a team.

7. Recommendation

Phase 1 (start now): Cloudflare Workers AI - ~Rs 500/month, private, zero setup. Best value; covers ~20 users.

Phase 2 (if speed matters / more users): DeepInfra (~Rs 1,800/mo) or RunPod Serverless RTX 4090 + vLLM (~Rs 1,600-2,000/mo) - fast, private.

Phase 3 (top quality + always-instant for a large team): rent RTX 5090 32GB (~Rs 27,000/mo).

Avoid for company data: OpenRouter free models and Gemini free tier (they train on our data); per-request premium models (~Rs 1 lakh/mo at our scale); CPU servers like DigitalOcean/Hetzner (cannot run the AI); Groq for a team (per-minute limit).

Summary: Start on Cloudflare Workers AI (~Rs 500/month, ~Rs 30 per user) - private, no setup. Move to DeepInfra or a rented RTX 4090 with vLLM (~Rs 1,600-2,000/month) only when speed or user count grows. Avoid OpenRouter free / Gemini free (they train on our data) and CPU-only servers (cannot run the AI).

Rates web-verified Jun 2026: OpenAI GPT-4o-mini \$0.15/\$0.60, GPT-4o \$2.50/\$10; Claude Haiku 4.5 \$1/\$5; DeepInfra Qwen3-30B \$0.20/\$0.80; Gemini 2.0 Flash \$0.10/\$0.40. Cloudflare/RunPod/Vast.ai from their live pages. Actual cost varies with real usage.